

# **CSCI 350 Project 4**

## **Part 2**

**Kivilcim Cumbul**

# **Task 2 cont'd: Page Fault method**

# Page Fault method in vm.c

1- Get cr2 address

2- Check if the address found by rcr2() method is not 0.

3- Find page table entry (PTE) → **hint: use walkpgdir method**

4-If pte is not shared → give panic error

5-If pte is not present → give panic error

6-Find physical address(pa) and 20-bit physical page number (ppn)

7- Check if the table is shared or not → **hint: you have to check it using counter in this stage**

8 - If shared → look next slides

9- If not shared → make table writable and not shared using flags

10- Reinstall TLB Entries → `lcr3(V2P(myproc()->pgdir));`

It must be shared to get a pagefault as described in the project 4 document

Use flags we discussed in part 1 to check shared and present

If it is not present, it is invalid.

Step by step explanation


# 1- Get cr2 address

There are 4 control registers: %cr0, %cr2, %cr3, and %cr4

We will get the address from cr2 register —> using rcr2() method

The method is defined in x86.h

```
static inline uint
rcr2(void)
{
    uint val;
    asm volatile("movl %%cr2,%0" : "=r" (val));
    return val;
}
```



## CR2

bit	label	description
0-31	pfla	page fault linear address

In xv6, cr2 register gives the virtual address that causes the page fault.

Check this for more info: [https://wiki.osdev.org/CPU\\_Registers\\_x86#CR2](https://wiki.osdev.org/CPU_Registers_x86#CR2)

## **2- Check if the address found by rcr2() method is not null.**

Check if the address is not 0.

This is important because we have to make sure the address we found is not null.

If it is null we have to throw null pointer exception or indicate it is null pointer using panic

Terminate the current process, and exit

### 3- Find page table entry (PTE)

```
// Return the address of the PTE in page table pgdir
// that corresponds to virtual address va. If alloc!=0,
// create any required page table pages.
static pte_t *
walkpgdir(pte_t *pgdir, const void *va, int alloc)
{
    pde_t *pde;
    pte_t *pgtab;

    pde = &pgdir[PDX(va)];
    if(*pde & PTE_P){
        pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
    } else {
        if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
            return 0;
        // Make sure all those PTE_P bits are zero.
        memset(pgtab, 0, PGSIZE);
        // The permissions here are overly generous, but they can
        // be further restricted by the permissions in the page table
        // entries, if necessary.
        *pde = V2P(pgtab) | PTE_P | PTE_W | PTE_U;
    }
    return &pgtab[PTX(va)];
}
```

Use walkpgdir method

## **4-Check if pte is shared or not**

If pte is shared, that means we are good to go, because 2 processes or more share the same page table entry which indicates a pagefault can be thrown.

If pte is not shared, the page fault cannot be thrown in the first place, so we should give an error message

To check, use the shared flag you created in the first part of project 4.

Use panic to throw error.



## 4-Check if pte is present or not

If pte is present, we can continue.

But if pte is not present, the page fault cannot be thrown in the first place, so we should give an error message.

To check, use the PTE\_P flag you learned in first part of the project

Use panic to throw error.

## 4-Find physical address(pa) and 20-bit physical page number (ppn)

Physical address (pa) is an address that the processor chip sends to main memory.

You can use PTE\_ADDR() method in mmu.h to find the pa

```
// Address in page table or page directory entry  
#define PTE_ADDR(pte) ((uint)(pte) & ~0xFFF)
```

Each PTE contains a 20-bit physical page number (PPN) and some flags. The paging hardware translates a virtual address by using its top 20 bits to index into the page table to find a PTE, and replacing the address's top 20 bits with the PPN in the PTE. The paging hardware copies the low 12 bits unchanged from the virtual to the translated physical address.

You can find ppn using shifting and/or masking

This step can be skipped in some implementations — depending on how you implemented the counter.

## **7– Check if the table is shared or not**

This time, we will use counter we created earlier to check if there are still process that shared this table.

Everyone can implement the counter differently, so there is no specific way to do this part

pa and/or ppn can be used here

## **8- If shared**

- 8.1 - create a new page table
- 8.2 - copy contents of page table we have
- 8.3 - change flags (user, present, shared flags)
- 8.4 - insert the new physical page num
- 8.5 - Decrease the counter

## 8.1 - create a new page table

## 8.2 - copy contents of page table we have

Hint : you might want to check these methods

### kalloc() method in kalloc.c → create

```
// Allocate one 4096-byte page of physical memory.
// Returns a pointer that the kernel can use.
// Returns 0 if the memory cannot be allocated.
char*
kalloc(void)
{
    struct run *r;

    if(kmem.use_lock)
        acquire(&kmem.lock);
    r = kmem.freelist;
    if(r)
        kmem.freelist = r->next;
    if(kmem.use_lock)
        release(&kmem.lock);
    return (char*)r;
}
```

### memmove() method in string.c → copy

```
void*
memmove(void *dst, const void *src, uint n)
{
    const char *s;
    char *d;

    s = src;
    d = dst;
    if(s < d && s + n > d){
        s += n;
        d += n;
        while(n-- > 0)
            *--d = *--s;
    } else
        while(n-- > 0)
            *d++ = *s++;

    return dst;
}
```

## 8.3 – change flags (user, writable, shared flags)

You have to do this stage because after creating a new page table and copying contents, the process will have a new and its own page table, this means:

-It is not shared

-It is writable

-It is user and present

```
// Page table/directory entry flags.  
#define PTE_P          0x001    // Present  
#define PTE_W          0x002    // Writeable  
#define PTE_U          0x004    // User  
#define PTE_PS         0x080    // Page Size
```

## 8.4 – insert the new physical page num

We need to set first 20 bits of the pte to the physical address of the new page table

This means we have to use either one of these methods: P2V or V2P. Choose the proper one.

```
11  #define V2P(a) (((uint) (a)) - KERNBASE)
12  #define P2V(a) ((void *)(((char *) (a)) + KERNBASE))
```

## 8.5 – Decrease the counter

Now, because we throw a pagefault interrupt, the tables of the child and parent got separated:

This means, we have to decrease the counter which we implemented earlier to store how many processes are sharing the parent's page table.

**The counter implementation can be different for every student**, so there is no specific way to decrement it.



## 9– If not shared:

Use flags to make the table not shared and writable

```
// Page table/directory entry flags.  
#define PTE_P      0x001    // Present  
#define PTE_W      0x002    // Writeable  
#define PTE_U      0x004    // User  
#define PTE_PS     0x080    // Page Size
```

And the flag you created

**This is the end of the pagefault method discussion,**  
but don't forget to reinstall TLB Entries!

# **Task 2 cont'd: changes in proc.c**

Now, instead of copyuvm() method, you have to call cow() method while **forking**.

## **Task 2 cont'd: changes in defs.h**

**Add definitions** for `cow()` and `pagefault()` methods we created in `vm.c`  
**to `defs.h`**

# **Task 2 cont'd: changes in trap.c**

# Add pagefault as a new type of trap in trap.c —> define new case

```
1 // x86 trap and interrupt constants.
2
3 // Processor-defined:
4 #define T_DIVIDE      0      // divide error
5 #define T_DEBUG       1      // debug exception
6 #define T_NMI         2      // non-maskable interrupt
7 #define T_BRKPT       3      // breakpoint
8 #define T_OFLOW       4      // overflow
9 #define T_BOUND       5      // bounds check
10 #define T_ILLOP       6      // illegal opcode
11 #define T_DEVICE       7      // device not available
12 #define T_DBLFLT      8      // double fault
13 // #define T_COPROC    9      // reserved (not used since 486)
14 #define T_TSS         10     // invalid task switch segment
15 #define T_SEGNP       11     // segment not present
16 #define T_STACK       12     // stack exception
17 #define T_GPFLT       13     // general protection fault
18 #define T_PGFLT       14     // page fault
19 // #define T_RES       15     // reserved
20 #define T_FPERR       16     // floating point error
21 #define T_ALIGN       17     // alignment check
22 #define T_MCHK        18     // machine check
23 #define T_SIMDERR     19     // SIMD floating point error
24
25 // These are arbitrarily chosen, but with care not to overlap
26 // processor defined exceptions or interrupt vectors.
27 #define T_SYSCALL      64     // system call
28 #define T_DEFAULT     500     // catchall
```

In traps.h —> T\_PGFLT(pagefault) flag is already defined:

When this case occurs pagefault method will be raised



# Task 3: testcow.c

To make sure that COW optimization implemented, use the included file `testcow.c`. Make sure your output matches the output below. Inspect the contents of `testcow.c` and explain what the role of `malloc()` and `fork()` are in the test (What are the functions doing specifically, what is the code doing at a high level and why does the output below show us COW is working correctly?). A brief description (3-5 sentences) should be sufficient. Include this explanation with your DESIGNDOC.

The test can be **different for each student** but make sure to have these:

`create values in different pages using malloc`

`Print parent info using procdump()`

`Fork to create new child processes`

`procdump()` to print info of child

`Update something in child process`  **Find it out!**

`procdump()` to print info of child

**That's all for project 4!**

Don't forget to add explanations about testcow.c to your DESIGNDOC.

**Stay safe !!!**