



# CSCI 350 Project 3

## Summer 2020

*MFQ Scheduler*

Kivilcim Cumbul



# Objectives

1. Implement a basic non-preemptive MFQ (Multilevel Feedback Queue)
  - ▶ Changes in `proc.c`, `proc.h`, `p_stat.h`, `trap.c`
2. Add a new system call to get information about a running process
  - ▶ `getpinfo(pid)`
3. Design tests to demonstrate the correctness of your scheduler.
  - ▶ `test1.c`, `test2.c`, `test3.c`
4. Create timeline graphs



# Before Starting...

Make sure you choose 1 CPU in your VM

And in your Makefile, replace CPUS := 2 with CPUS := 1



# Part 1

In this project, you'll implement a simplified Multi-level Feedback Queue (MFQ) scheduler in xv6.

- ▶ proc.c, proc.h, pstat.h, trap.c

Chapter 5 of the xv6 manual.

The default and only scheduling policy in xv6 is round-robin, and we will change it!

Note: I will explain more high level this time because everyone can implement this differently



# pstat.h

```
#ifndef _PSTAT_H_
#define _PSTAT_H_

/* NSCHEDSTATS is the number of sched_stat_t slots per process. The scheduler fills in the slots
as it schedules processes to record information about scheduling */

#define NSCHEDSTATS 1500

/*
 * responsible for recording the scheduling state per process at a particular tick
 * e.g. a process can have an array of sched_stat_t's, with each of them holding the info
 * of a scheduling round of the process
 */
struct sched_stat_t
{
    int start_tick; //the number of ticks when this process is scheduled
    int duration; //number of ticks the process is running before it gives up the CPU
    int priority; //the priority of the process when it's scheduled

    //you may add more fields for debugging purposes
};

#endif
```

Record scheduling stats  
And use them again

It fills array for each process  
and records info of its scheduling

DONT FORGET in proc.c — #include "pstat.h"



# proc.h

## Inside proc structure

```
int times[3]           // number of times each process was scheduled at each of 3  
                      // priority queues  
  
int ticks[3]           // number of ticks each process used the last time it was  
                      // scheduled in each priority queue  
                      // cannot be greater than the time-slice for each queue  
  
uint wait_time;        // number of ticks each RUNNABLE process waited in the lowest  
                      // priority queue
```



# wait...

```
struct sched_stat_t
{
    int start_tick; //the number of ticks when this process is scheduled
    int duration; //number of ticks the process is running before it gives up the CPU
    int priority; //the priority of the process when it's scheduled

    //you may add more fields for debugging purposes
};

#endif
```



```
int times[3] // number of times each process was scheduled at each of 3
              // priority queues

int ticks[3] // number of ticks each process used the last time it was
              // scheduled in each priority queue
              // cannot be greater than the time-slice for each queue

uint wait_time; // number of ticks each RUNNABLE process waited in the lowest
                 // priority queue
```

We can put these either proc struct  
or stat struct

The only important thing is  
recording!



# How do we call them?

In proc structure

```
proc->times[0]  
proc->times[1]  
proc->times[2]  
  
proc->ticks0]  
proc->ticks[1]  
proc->ticks[2]
```

```
proc->wait_time
```

In sched\_stat\_t

```
proc->sched_stat_t[pid].start_tick  
proc->sched_stat_t[pid].duration  
proc->sched_stat_t[pid].priority
```



# proc.h additional variables

You might define as much variables as you want into proc or sched\_stats\_t.  
Some examples for proc:

```
int ticks;           //number of timer ticks the process has run for  
  
int total_ticks;    //total number of timer ticks the process has run for  
  
struct sched_stat_t sched_stats[NSCHEDSTATS]; // schedule stats for each tick  
  
int num_stats_used; // count to the end of the array
```

**Dont forget this!**

  
Because array is predefined with  
the value we defined we have keep  
track where is the real end



# Let's start scheduler! — Part 1

It is much easier to deal with fixed-sized arrays in xv6 than linked-lists. For simplicity, we recommend that you use arrays to represent each priority level (queue).

```
struct proc* q0[NPROC]; → Level 1 — First → 1 tick  
struct proc* q1[NPROC]; → Level 2 → 2 ticks  
struct proc* q2[NPROC]; → Level 3 — Last → 8 ticks
```

## proc.h

```
extern struct proc* q0[NPROC];  
  
extern struct proc* q1[NPROC];  
  
extern struct proc* q2[NPROC];
```

## proc.c

```
struct proc* q0[NPROC];  
  
struct proc* q1[NPROC];  
  
struct proc* q2[NPROC];
```



# Missing something?

```
struct proc* q0[NPROC]; → Level 1 — First  
struct proc* q1[NPROC]; → Level 2  
struct proc* q2[NPROC]; → Level 3 — Last
```

Can we use something else?

COUNTERS!!

Sure!  
Tail, Front, Capacity etc.

proc.h

```
extern int c0;  
  
extern int c1;  
  
extern int c2;
```

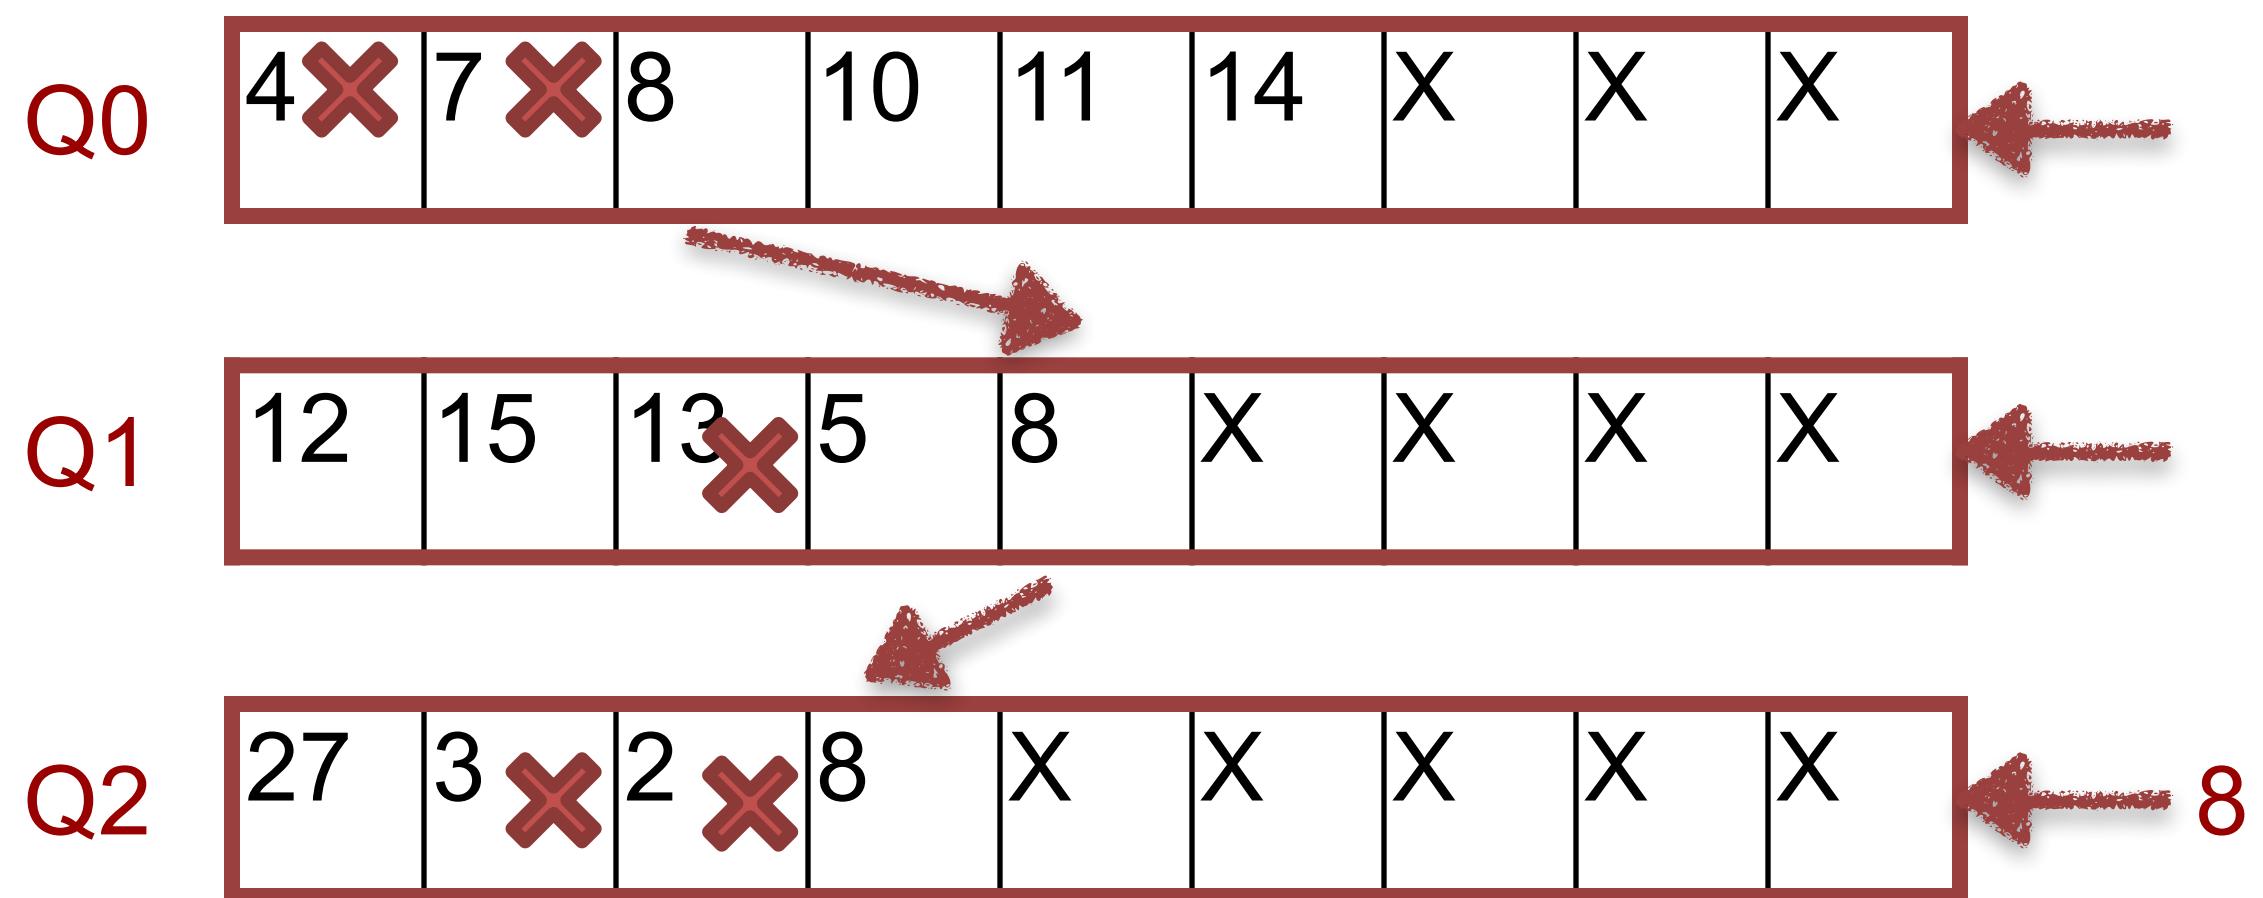
proc.c

```
int c0=-1;  
  
int c1=-1;  
  
int c2=-1;
```



# Visualization

Choose whoever is runnable (8 in this case)



- 1 - Numbered from 0 (highest) down to 2 (lowest).
- 2 - When ticks, highest priority **RUNNABLE** process is scheduled to run.
- 3 - If more than one process on the same level, then use **round robin fashion**.
- 4 - Q0 → 1 timer tick. Q1 → 2 timer ticks Q2 → 8 timer ticks.
- 5- Can get out before CPU increase the tick → **CHEAT**
- 6- When a new process arrives, it **should start at priority 0** (highest priority).  
Place this process at the **end of the highest priority queue**
- 7 - At priorities 0, and 1, after a process consumes its time-slice it should be downgraded one priority level. Whenever a process is moved to a **lower priority** level, it should be **placed at the end of the queue**.
- 8 - If a process wakes up after **voluntarily giving up the CPU** it stays at the same priority level. Place this process at the end of its queue; it should not preempt a process with the same priority. → **CHEAT**
- 9- Your scheduler should never **preempt a lower priority process** if a higher priority process is available to run.

In proc.c



# In allocproc()

## Allocate Process

Look in the process table for an UNUSED proc.

If found, change state to EMBRYO and initialize state required to run in the kernel.

Otherwise return 0.

This means we need to check our queues and if there is a unused etc. remove that process

After found -> Remove from queue:

```
if (p->pid > 0):
    for (i=0; (i<c0); i++)
        if p == q0[i]
            Delete q0[i]
            Shift all left
```

```
for (i=0; (i<c1); i++) →
...
for (i=0; (i<c2); i++) →
```

Do the same thing for  
q1 and q2

```
allocproc(void)
{
    struct proc *p;
    char *sp;

    acquire(&ptable.lock);

    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
        if(p->state == UNUSED)
            goto found;

    release(&ptable.lock);
    return 0;

found:
```

In proc.c

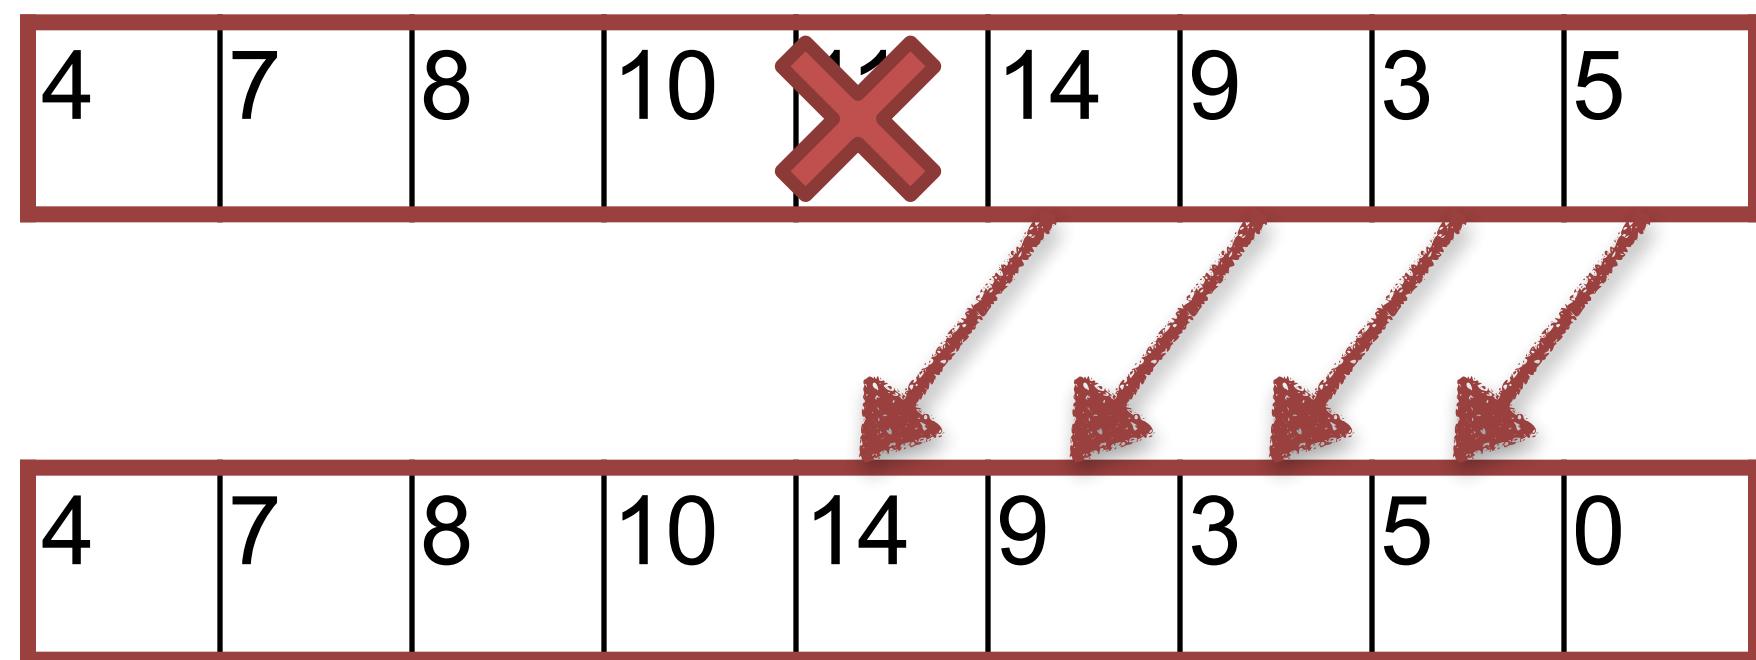


# Shift all left?

Lets say we are looking for pid = 11

4	7	8	10	11	14	9	3	5
---	---	---	----	----	----	---	---	---

Not need to delete just shifting is enough



Last element?

$q0[c0]=0 ; (c0, c1, c2)$

Also counter  $\rightarrow$  the tail of the array changes:

$counter-- ; (c0, c1, c2)$

In proc.c



# In allocproc()

This might not be necessary if proc you  
Did not define any variables to initiate in proc

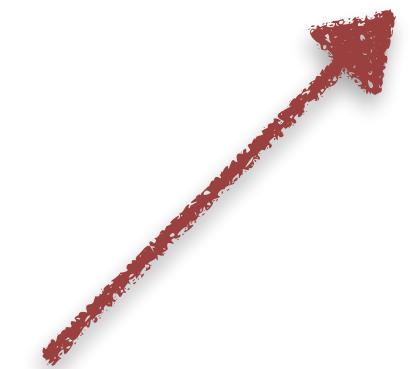
After searching all and shifting:

Initiate proc properties:

Priority, total\_ticks etc -> whatever is necessary

c0++;

p = q0[c0] -> Add p into first priority queue



In proc.c



# In userinit()

## User initialize

Everything starts with userinit system call

Reset schedule stats because we just start the first process here

From, xv6 manual pg 24

### Code: creating the first process

Now the kernel runs within its own address space, we look at how the kernel creates user-level processes and ensures strong isolation between the kernel and user-level processes, and between processes themselves.

After main initializes several devices and subsystems, it creates the first process by calling `userinit` (1239). Userinit's first action is to call `allocproc`. The job of al-

Reset here

```
// Set up first user process.  
void  
userinit(void)  
{  
    struct proc *p;  
    extern char _binary_initcode_start[], _binary_initcode_size[];  
  
    p = allocproc();  
  
    initproc = p;  
    if((p->pgdir = setupkvm()) == 0)  
        panic("userinit: out of memory?");  
    inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);  
    p->sz = PGSIZE;  
    memset(p->tf, 0, sizeof(*p->tf));  
    p->tf->cs = (SEG_UCODE << 3) | DPL_USER;  
    p->tf->ds = (SEG_UDATA << 3) | DPL_USER;  
    p->tf->es = p->tf->ds;  
    p->tf->ss = p->tf->ds;  
    p->tf->eflags = FL_IF;  
    p->tf->esp = PGSIZE;  
    p->tf->eip = 0; // beginning of initcode.S
```

In proc.c

# scheduler()

## Userinit calls it at the beginning

Now that the first process's state is prepared, it is time to run it. After `main` calls `userinit`, `mpmain` calls `scheduler` to start running processes (1267). `Scheduler` (2708) looks for a process with `p->state` set to `RUNNABLE`, and there's only one: `initproc`. It sets the per-cpu variable `proc` to the process it found and calls `switchuvm` to tell the hardware to start using the target process's page table (1868). Changing page tables

From, xv6 manual pg 26

`scheduler` now sets `p->state` to `RUNNING` and calls `swtch` (2958) to perform a context switch to the target process's kernel thread. `swtch` saves the current registers and loads the saved registers of the target kernel thread (`proc->context`) into the x86 hardware registers, including the stack pointer and instruction pointer. The current context is not a process but rather a special per-cpu scheduler context, so `scheduler` tells `swtch` to save the current hardware registers in per-cpu storage (`cpu->scheduler`) rather than in any process's kernel thread context. We'll examine `swtch` in more detail in Chapter 5. The final `ret` instruction (2977) pops the target process's `%eip` from the stack, finishing the context switch. Now the processor is running on the kernel stack of process `p`.

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

            // Switch to chosen process.  It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

In proc.c



# swtch()

Per-CPU process scheduler. Each CPU calls scheduler() after setting itself up.

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();
        }
        // Process is done running for now.
        // It should have changed its p->state before coming back.
        c->proc = 0;
    }
    release(&ptable.lock);
}
```

Scheduler never returns. It loops, doing:

- choose a process to run
- swtch to start running that process
- eventually that process transfers control via swtch back to the scheduler.

```
void
sleep(void *chan, struct spinlock *lk)
{
    struct proc *p = myproc();

    if(p == 0)
        panic("sleep");

    if(lk == 0)
        panic("sleep without lk");

    if(lk != &ptable.lock){ //DOC: sleeplock0
        acquire(&ptable.lock); //DOC: sleeplock1
        release(lk);
    }
    // Go to sleep.
    p->chan = chan;
    p->state = SLEEPING;
    sched();
}

// Give up the CPU for one scheduling round.
void
yield(void)
{
    acquire(&ptable.lock); //DOC: yieldlock
    myproc()->state = RUNNABLE;
    sched();
    release(&ptable.lock);
}
```

exit() —> sched()

```
void
sched(void)
{
    int intena;
    struct proc *p = myproc();

    if(!holding(&ptable.lock))
        panic("sched ptable.lock");
    if(mycpu()->ncli != 1)
        panic("sched locks");
    if(p->state == RUNNING)
        panic("sched running");
    if(readeflags()&FL_IF)
        panic("sched interruptible");
    intena = mycpu()->intena;
    swtch(&p->context, mycpu()->scheduler);
    mycpu()->intena = intena;
}
```

In proc.c

# switchkvm/switchuvvm

From, xv6 manual:

The **u** in switchuvvm stands for User. The **k** in switchkvm stands for Kernel.

User mode can call system call

When changing protection levels from user to kernel mode, the kernel shouldn't use the stack of the user process, because it may not be valid. The user process may be malicious or contain an error that causes the user %esp to contain an address that is not part of the process's user memory. Xv6 programs the x86 hardware to perform a stack switch on a trap by setting up a task segment descriptor through which the hardware loads a stack segment selector and a new value for %esp. The function `switchuvvm` (1873) stores the address of the top of the kernel stack of the user process into the task segment descriptor.

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p; Add process to cpu
            switchuvvm(p);
            p->state = RUNNING;
        }

        swtch(&(c->scheduler), p->context);
        switchkvm();
    }

    // Process is done running for now.
    // It should have changed its p->state before coming back.
    c->proc = 0;
}

release(&ptable.lock);

}
```

Before running

After running

In proc.c



## How to modify?

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

Same

For loop finds next available running thread  
But it does not work like that in MFQ

If q0 is not empty:  
Loop to find RUNNABLE process in the q0 array  
if q0[i] ->state != RUNNABLE  
Continue  
p = q0[i] **Found first runnable**  
c->proc = p  
switchuvm(p)  
p->state = RUNNING

WRITE STATS → start\_tick and priority

swtch(&(c->scheduler), p->context);  
switchkvm()

start\_tick = ticks

Came back here

Duration = ticks - start  
p->num\_stats\_used++;  
p->times[0]++;  
p->ticks[0] = duration;  
Change total\_ticks in proc if you have

In proc.c



# How to think about stats?

Arbitrary program in any language  
How to time?

Start time = x

Do some work

Go to different functions

Come back

End time = y

Duration = y-x

```
#ifndef _PSTAT_H_
#define _PSTAT_H_

/* NSCHEDSTATS is the number of sched_stat_t slots per process. The scheduler
fills in the slots
as it schedules processes to record information about scheduling */

#define NSCHEDSTATS 1500
struct sched_stat_t
{
    int start_tick; //the number of ticks when this process is scheduled
    int duration;   //number of ticks the process is running before it gives up
    the CPU
    int priority;   //the priority of the process when it's scheduled
                    //you may add more fields for debugging purposes
};

#endif
```

In proc.c

# tick, proc->tick, proc->tick[0]



tick comes from trap.c

proc->tick — extra : (increase in trap.c)  
number of timer ticks the process has run for

proc->tick[] array:  
number of ticks each process used the last time it was  
scheduled in each priority queue

proc->times[] array:  
number of times each process was scheduled at each of 3

```
*****
name= CPUintensive, pid= 5
wait time= 0
ticks= {1, 2, 8}
times= {1, 1, 3}
*****
start=1, duration=1, priority=0
start=2, duration=0, priority=1
start=2, duration=2, priority=1
start=4, duration=8, priority=2
start=12, duration=8, priority=2
```

If q0 is not empty:  
Loop to find RUNNABLE process in the q0 array  
if q0[i] ->state != RUNNABLE  
Continue

p = q0[i] **Found first runnable**  
c->proc = p  
switchuvm(p)  
p->state = RUNNING

WRITE STATS → start\_tick and priority **start\_tick + ticks**

swtch(&(c->scheduler), p->context);  
switchkvm()

Duration = **ticks** - start  
p->num stats used++;  
p->**times[0]**++;  
p->**ticks[0]** = duration;  
Change **total ticks** if you have

In proc.c

# How to modify? cont'd



```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

## Second half

If q0 is not empty:

... Previous slide

if(p->ticks >= 1)

p->ticks = 0;

copy proc to lower priority queue

increase c1 count

... other changes necessary

q1[c1] = p; // put at the end of the q1

q0[i] = 0; // delete proc from q0

Shift all left from i **Dont forget counter-- and last element of array**

boost()

c->proc = 0; **Reset cpu**

In proc.c

# How to modify?



# Q1

```
If q1 is not empty:  
    Loop to find RUNNABLE process in the q1 array  
        if q1[i] ->state != RUNNABLE  
            Continue  
        p = q1[i]  
        c->proc = p  
        switchuvm(p)  
        p->state = RUNNING  
  
        WRITE STATS -> start_tick and priority start_tick = ticks  
  
        swtch(&(c->scheduler), p->context);  
        switchkvm()  
  
        Duration = ticks - start  
        p->num_stats_used++;  
        p->times[1]++;  
        p->ticks[1] = duration;  
        Change total_ticks in proc if you have
```

In proc.c



# Q1

## Second half

If q1 is not empty:

... Previous slide

```
if(p->ticks >= 2)
    p->ticks = 0;
    copy proc to lower priority queue
    increase c2 count
    ... other changes necessary
    q2[c2] = p; // put at the end of the q2
    q1[i] = 0; // delete proc from q1
    Shift all left from i Dont forget counter-- and last element of array
```

```
boost()
c->proc = 0; Reset cpu
```

In proc.c

# How to modify?



# Q2

```
If q2 is not empty:  
    Loop to find RUNNABLE process in the q2 array  
        if q2[i] ->state != RUNNABLE  
            Continue  
        p = q2[i]  
        c->proc = p  
        switchuvm(p)  
        p->state = RUNNING  
  
        WRITE STATS -> start_tick and priority start_tick = ticks  
  
        swtch(&(c->scheduler), p->context);  
        switchkvm()  
  
        Duration = ticks - start  
        p->num_stats_used++;  
        p->times[2]++;  
        p->ticks[2] = duration;  
        Change total_ticks in proc if you have
```

In proc.c



# Q2

## Second half

If q1 is not empty:

... Previous slide

```
boost()
if(p->ticks >=8)
    p->ticks = 0;
```

increase c2 count

**move process to end of its own queue**

**Shift left**

**q2[c2] = p**

c->proc = 0; **Reset cpu**

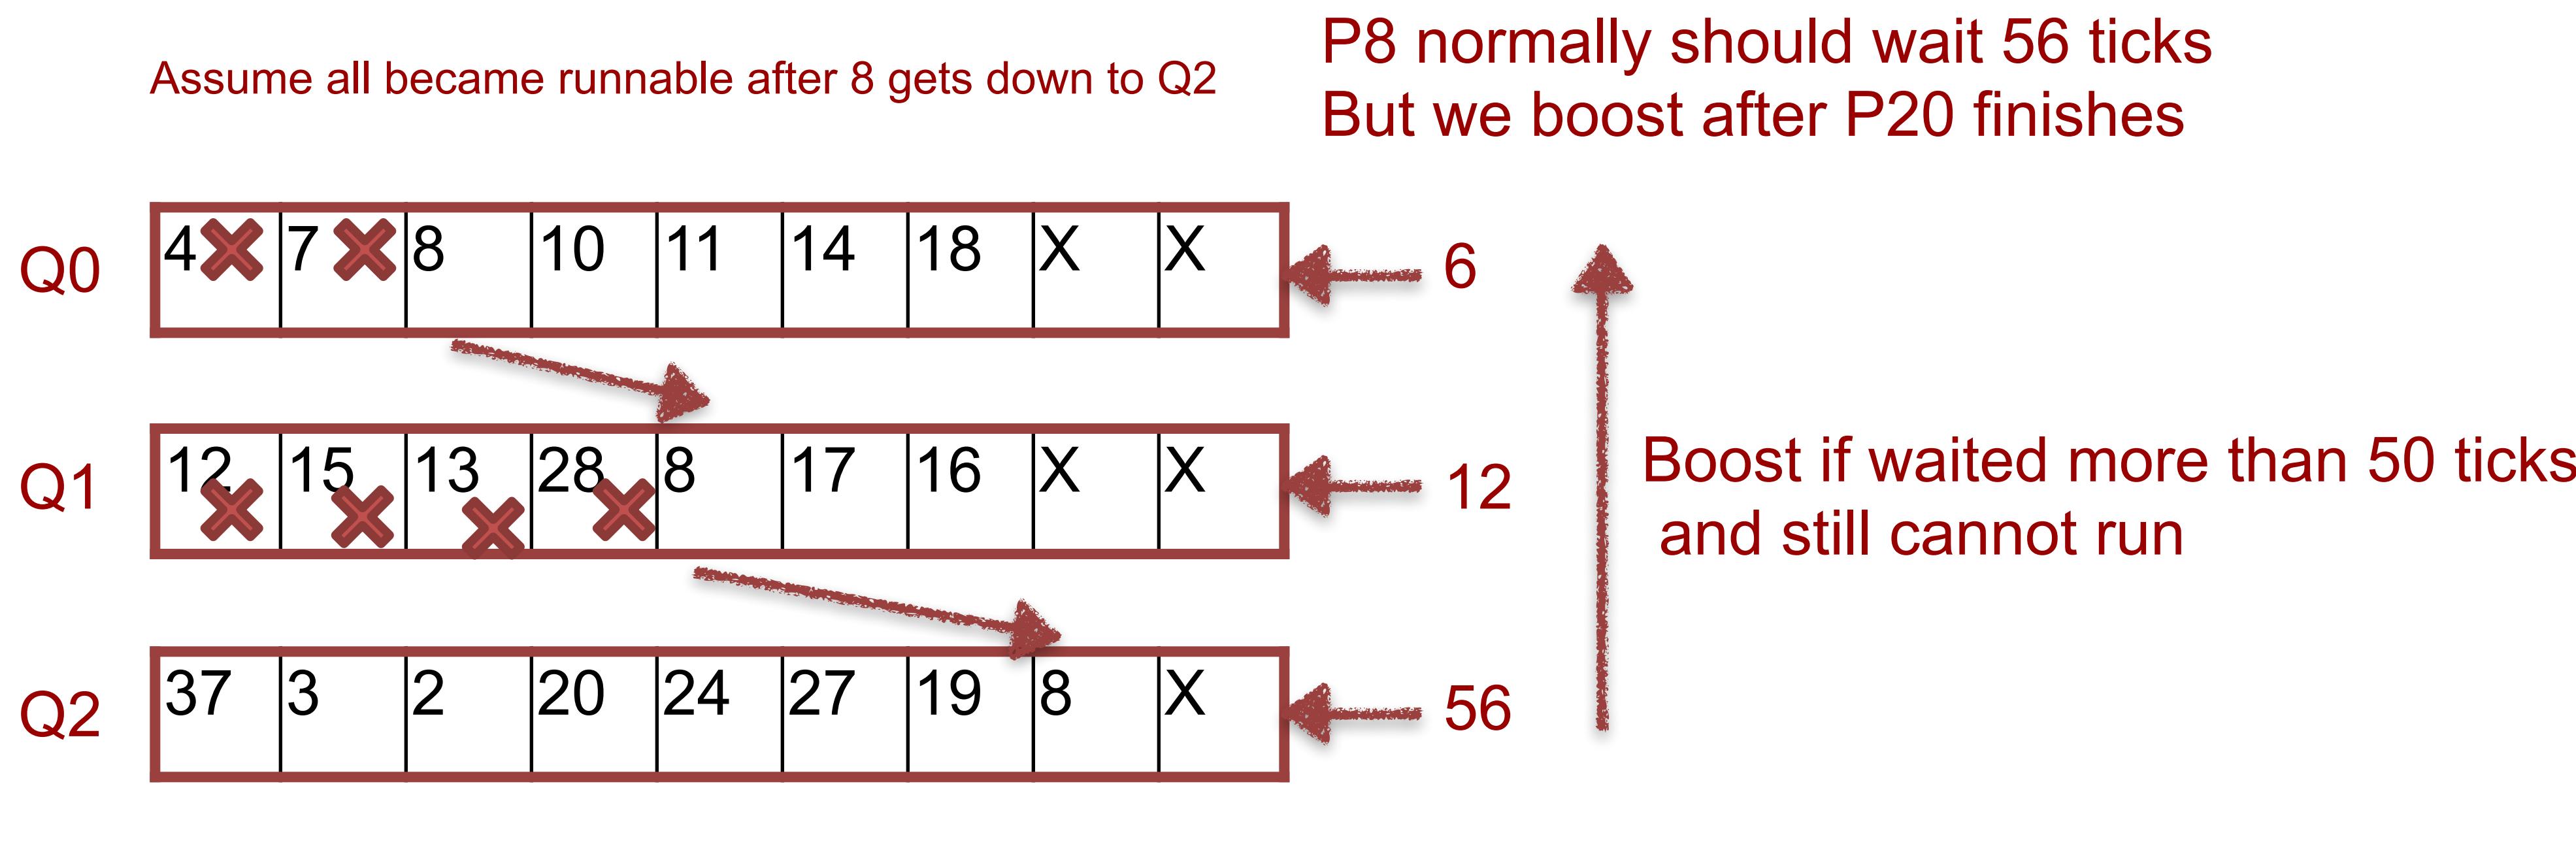
In proc.c



# boost()

Task 1-10: You need to implement the priority boosting mechanism which will be used to increase a priority of a process that has not been scheduled in a while. The goal is to **avoid starvation**, which happens when a process never receives CPU time because higher-priority processes keep arriving.

After a **RUNNABLE** process has been waiting in the lowest priority queue for **50 ticks or more**, move the process to the **end of the highest priority queue**. This method of priority boosting is called **aging**.





# boost()

boost (??) :

Increase wait time of all RUNNABLE processes in **q2** according to last process' duration

If any of the process' wait time passes 50  
boost it into the end of first priority queue

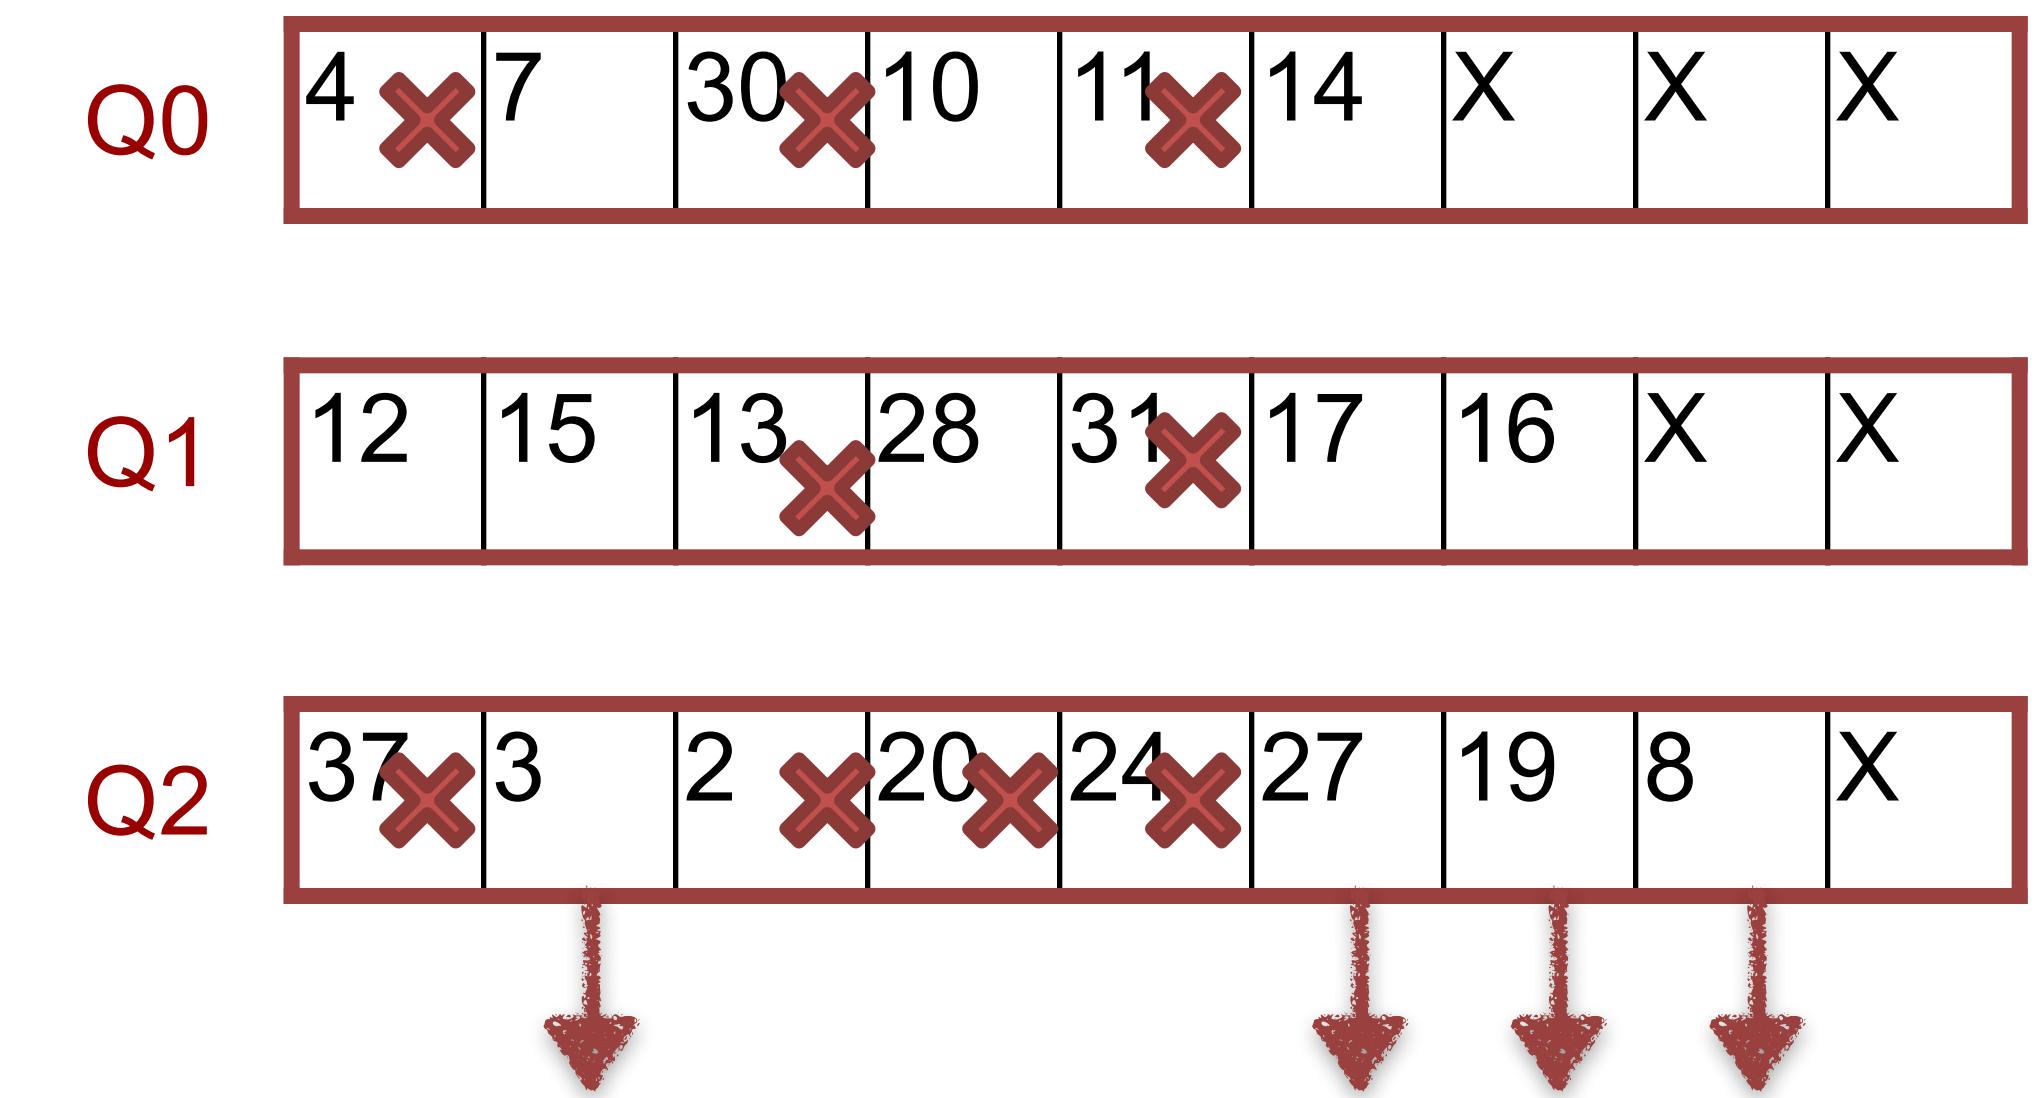
p = p\_to\_boost

Change priority of p

c0++;

q0[c0] = p

Delete from q2 and shift left



When 7 works +1

When 10 works +1

When 12 works +2

When 15 works +2

When 3 works +8

In proc.c



# getpinfo()

- Takes process id inside —> getpinfo(int pid)
- This is a system call so make sure to define it
- Define it in proc.h

syscall.h  
syscall.c  
user.h  
usys.S  
sysproc.c

int getpinfo(int pid);

```
*****  
name= CPUintensive, pid= 5  
wait time= 0  
ticks= {1, 2, 8}  
times= {1, 1, 3}  
*****  
start=1, duration=1, priority=0  
start=2, duration=0, priority=1  
start=2, duration=2, priority=1  
start=4, duration=8, priority=2  
start=12, duration=8, priority=2
```

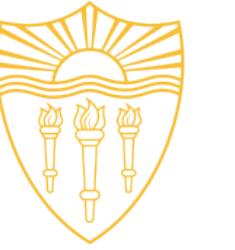
We will not grade ticks and times because it can be interpreted and implemented differently



# getpinfo()

```
getpinfo () :  
    Get lock  
    Get current process  
    With for loop through all processes in ptable  
        find pid  
    If not found  
        Release break  
    Cprintf  
  
    proc->name  
    proc->wait_time  
    proc->ticks[0],ticks[1],ticks[2]  
    proc->times[0],times[1],times[2]  
  
    // for each stat  
    For each stat until num_of_stats  
        Get/cprintf start_tick, duration and priority from each valid item in sched_stats  
    Release lock  
    Return 0
```

```
*****  
name= CPUintensive, pid= 5  
wait time= 0  
ticks= {1, 2, 8}  
times= {1, 1, 3}  
*****  
start=1, duration=1, priority=0  
start=2, duration=0, priority=1  
start=2, duration=2, priority=1  
start=4, duration=8, priority=2  
start=12, duration=8, priority=2
```



# Test 1

```
Fork()

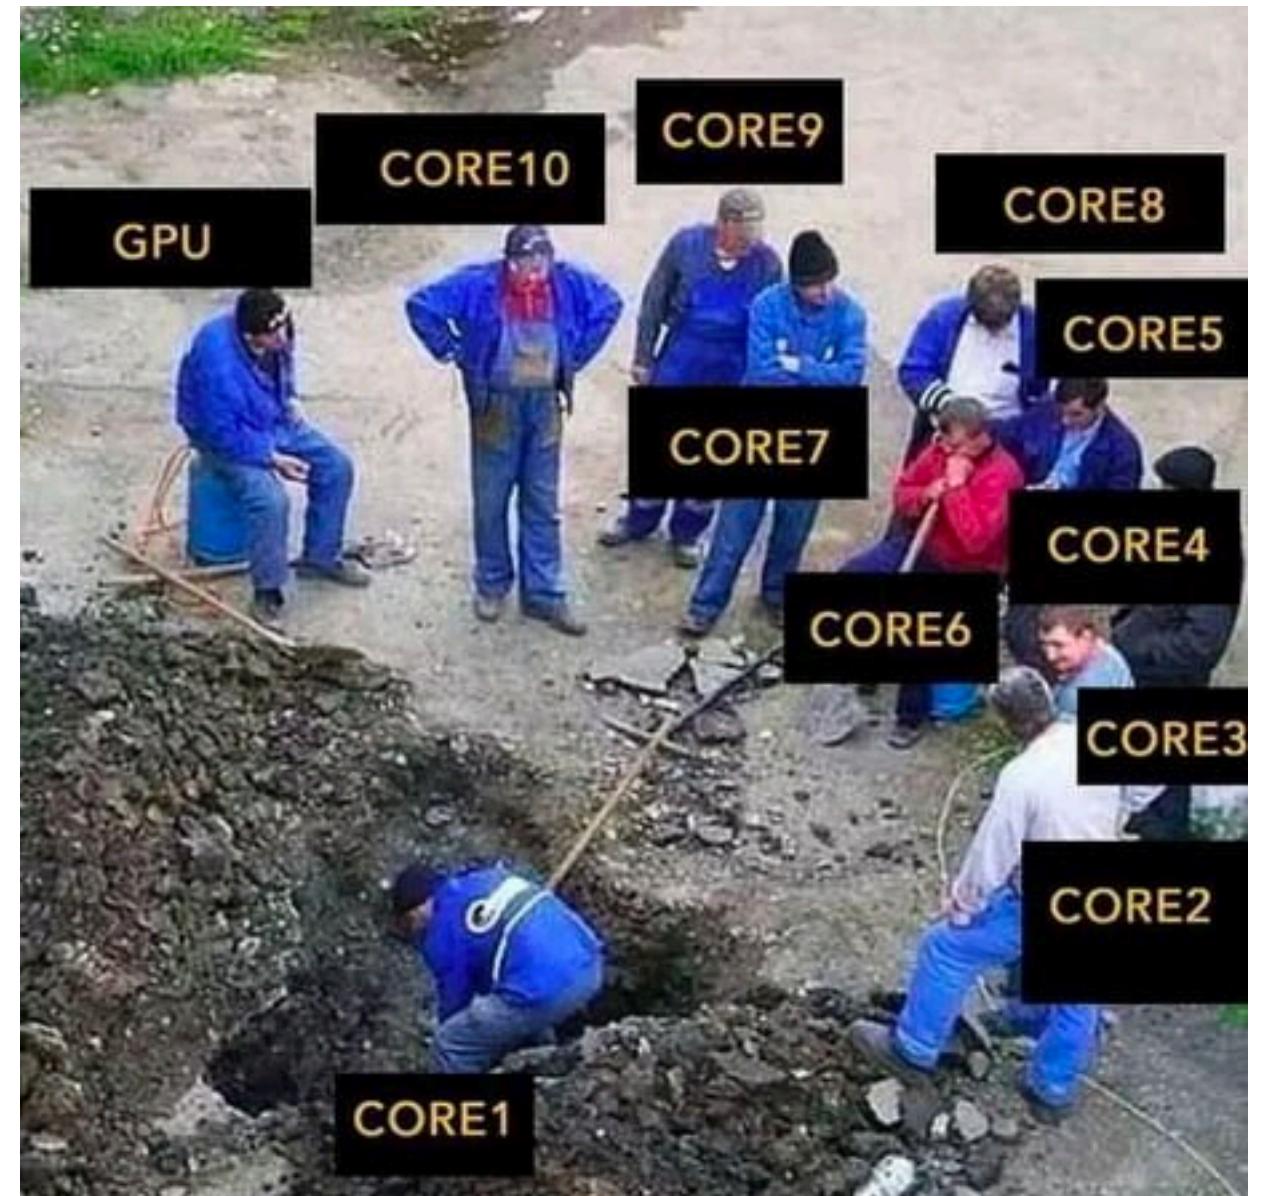
If (pid > 0)
    Do I/O instensive job
    wait();
    getpinfo(getpid());

Else:
    Do CPU intensive job
    exit()
```

## Hints:

Do not need to use input in I/O intensive job  
You can use output its input/output job!

Do not need to use recursion in CPU intensive job  
Basic calculations are enough in for loop





# Test 2

```
fork()  
In parent:  
    Fork multiple times – Like 10 times  
    Do CPU or I/O intensive job  
    getpinfo()
```

```
In child:  
    Do CPU or I/O intensive job in each task  
    getpinfo()
```

## Hints:

Make sure your CPU job is taking long time  
else you cannot see boost!





# Test 3

Hints:

Check slide 12 again

When you're thinking about going to sleep  
but now you can't sleep because you're  
thinking too much about going to sleep





# Graphing

Use excel or python - matplotlib

I might provide excel template later on Piazza

The link below is just an example it is not guaranteed that it works:

<https://github.com/joshmin98/auto-graph-project3/>

## Any questions?